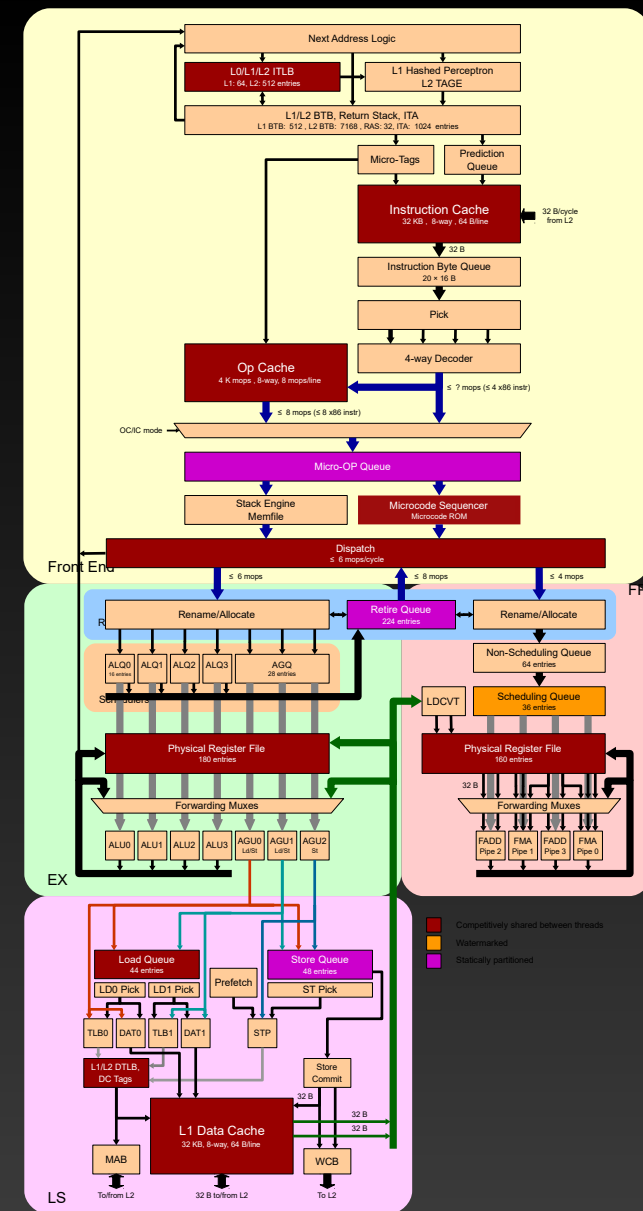


# Concurrent vs. parallel vs. distributed

- Concurrent
  - different parts of the program running in single system at the same time, can communicate
- Parallel
  - single computation divided to smaller and same parts, running concurrently
- Distributed
  - concurrent execution on different computers

# HW parallelism

- Registers
  - SSE = 128 bits
  - parts of  $2^n$  bytes
- Instructions
  - more ALUs
  - Out Of Order (OOO) execution
- SMT = Simultaneous Multi-Threading
  - more threads (2-16) in single core
  - better ALU utilization
- SMP – Symmetric Multi-Processing
- Cluster – computers + FAST interconnect
- Grid, Cloud – computers + common network



# Shared data

- OK only if single process (thread) is changing data
- Locking – necessary if more than one writer
- Locks
  - mutex, critical section, synchronized methods, semaphor, ...

# Implicit vs. explicit parallelism


- Implicit
  - Automatic parallelisation by compiler
    - both instruction level and pieces of source code (usually for loops)
  - Precompiled libraries (OpenCV)
  - Implicitly parallel programming languages
    - LabView, Matlab ( 1:N )
  - No effort to splitting, comm, sync
  - Smaller control over runtime, smaller efficiency, overhead is hidden

# Implicit vs. explicit parallelism


- Explicit
  - Precise control of concurrency, sync, comm using compiler directives, function calls etc.
    - overhead is visible and controllable
  - Directives
    - code splitting
    - synchronisation
    - communication
  - Full control, higher efficiency possible
  - thread API, OpenMP, MPI

# Data vs Task Parallelism

- Data-parallel
  - data are distributed
  - thread code (nearly) same



- Task-parallel
  - code is distributed



# Splitting APP to threads

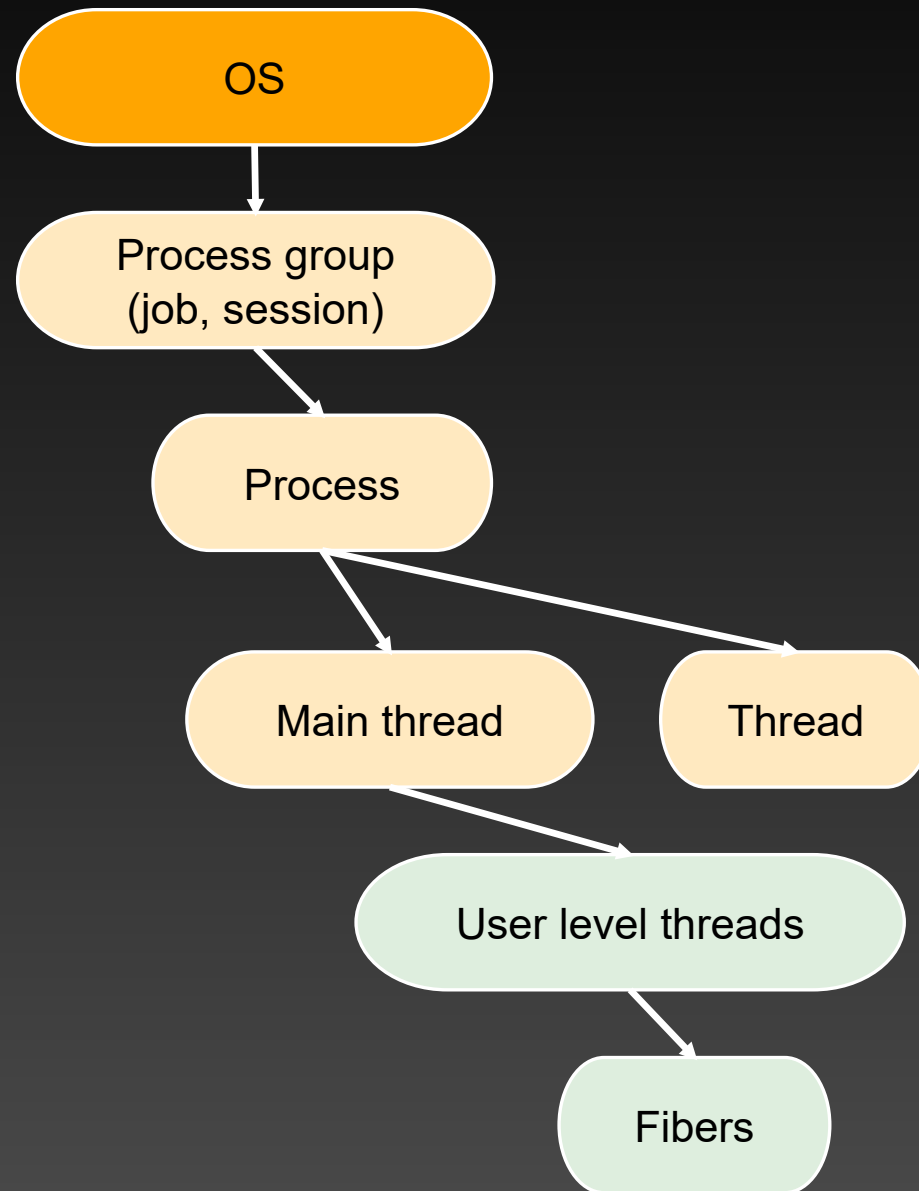
- Usually data and task parallelism
- Design patterns
  - Master/Slave + thread pool
    - master thread scatter and gather data + control others, does NOT compute itself
    - workers(slaves) threads usually created in advance
  - Equal threads
    - master also works
  - Pipeline
    - task parallel, each threads does different task
    - problems if one stage is slower

# Types of parallelism

- Fine grained
  - frequent comm and sync
  - small data blocks after shot execution
  - very latency sensitive
- Coarse grained
  - occasional communication
  - sync necessary, but not latency sensitive
- Embarrassingly parallel
  - completely independent tasks, zero comm
  - e.g. repeated run with different cmd line args

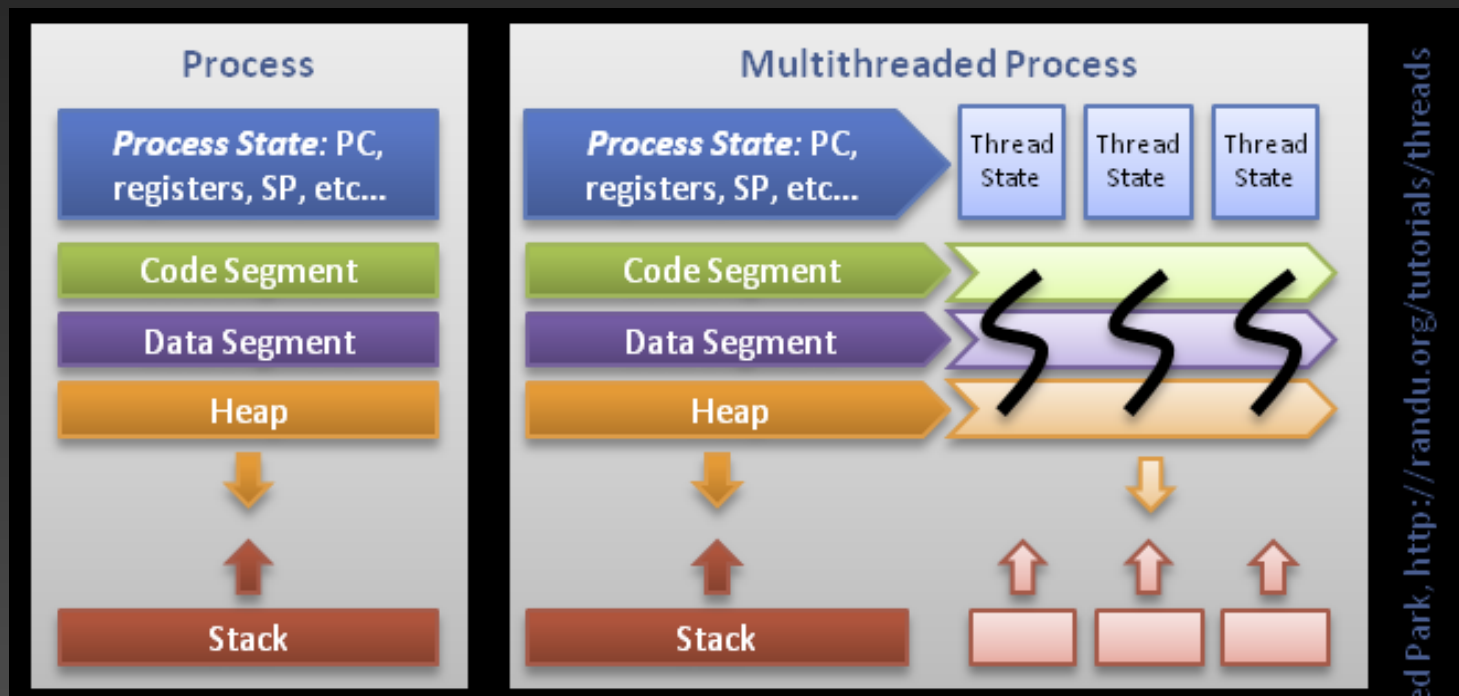


# Task Hierarchy



# Process vs. thread


- Process
  - Virtual memory, privileges, code, PID, priority
  - at least one thread
- Thread
  - memory is shared for all threads
  - thread local: only stack, registers, thread ID



# Synchronisation

- MUTEX = MUTual EXclusion
  - lock used for serialisation of thread access to resources
- Critical section
  - code between mutex locking and unlocking
  - guaranteed to be executed by only single thread at once
    - serial time
  - must be as small as possible
- Atomic operations
  - simple operations, guaranteed by hardware or library to be correct without explicit mutex (e.g. ++)

# Parallel vs. serial time



# Serial time impact

Two independent parts

**A** **B**

Original process



Make **B** 5x faster



Make **A** 2x faster



# Safety

- Dangerous operation considering parallel execution
  - uncontrolled access to globals (variables and heap)
  - saving function state to global variables
  - global resource (de)allocations (files, sockets, ... )
  - indirect access to data using pointers and references
  - visible side effects (modifications of volatile variables)
- Safe strategy
  - use only local variables (stack)
  - code depends only on function arguments, value passing
  - all functions and subfunctions are re-entrant

# Synchronisation primitives I

- Mutex – lock (locked vs. unlocked)
- Barrier
  - position in code, where execution of a thread is paused until all threads will arrive
- Join (fork-join)
  - gather results and exit status from all threads – join will terminate thread

# Synchronisation primitives II

- Conditional variable

- call wait() for variable → thread put to sleep
- HW watching for write into variable → wake up
- have to check for wake-up reason

```
lock( mutex_x );  
while ( not wake_me ) { sleep(cond_var, mutex_x); }  
unlock( mutex_x );
```

- Semaphore

- binary = mutex
- counted – set to N (resources), thread enters → --, thread exits → ++, on zero → wait



# Posix Threads

# POSIX vs. Win32 Threads

- Similar capabilities
  - win32 handle (=32bit int, for everything) vs. strong typing (each object has own data type)
  - POSIX officially only for C
- implementation of pthreads using win32
  - `#include <pthread.h>`
  - `#include <winpthreads.h>`
- <http://randu.org/tutorials/threads/>
- [http://locklessinc.com/articles/pthreads\\_on\\_windows](http://locklessinc.com/articles/pthreads_on_windows)

# Calls

- Over 60 API functions
- prefix for entity type, suffix for operation
  - pthread\_, pthread\_attr\_
  - pthread\_mutex\_, pthread\_mutexattr\_
  - pthread\_cond\_, pthread\_condattr\_
  - pthread\_key\_

# Attributes

- Properties of entity (thread, mutex, cond. variable) is set with special objects – attribute objects
- Some entity properties must be specified before entity creation
- Attribute object types
  - Thread: `pthread_attr_t`
  - Mutex: `pthread_mutexattr_t`
  - Conditional variable: `pthread_condattr_t`
- Creation and destruction
  - function `_init(...)` and `_destroy(...)` with prefix
  - parameter set to pointer to attribute object

# Thread creation

- Each program has one main thread, created by OS
- Other threads created explicitly
- Each thread can start more threads
- Thread is created by `pthread_create()`
- Thread is immediately ready to run
  - It can be started by OS scheduler before parent thread returns from `pthread_create()` function
  - All data necessary for thread must be prepared BEFORE calling `pthread_create()`

# Thread creation

- `int pthread_create( pthread_t *thread_handle,  
                  const pthread_attr_t *attribute,  
                  void * (*thread_function)(void *), void *arg);`
- `thread_handle` – thread descriptor
- `attribute` – pointer to structure with attributes of created thread (NULL for standard settings)
- `thread_function` – pointer to function to execute in thread
- `arg` – pointer to parameters of `thread_function`
- returns 0 if successful

# Thread properties

- Detached threads
  - Can not be joined with master by `pthread_join()`
  - Run on background, saves app resources
  - Standard thread properties are not always obvious
    - explicit setting recommended
    - `int pthread_detach ( pthread_t *thread_handle)`
    - `int pthread_attr_t setdetachstate(pthread_attr_t *attr, int detachstate)`
    - `int pthread_attr_t getdetachstate( pthread_attr_t *attr, int *detachstate)`

# Terminating thread

- Thread can be terminated
  - by calling `pthread_exit()` from inside
  - by ending parent thread execution by different call than `pthread_exit()` ( e.g. `exit()`, `abort()`, `return`, ...)
  - by cancelling using `pthread_cancel()`
  - by ending master thread other than `return` (`kill`, `exit`, `abort`...)
- `void pthread_exit (void *value)`
  - terminates thread execution
  - process resources (`fd`, `IPC`, `mutex`, ... ) created (opened) in thread are NOT closed (deallocated) – global resources
  - heap data referenced only from thread must be released before exit – memory leak (system will release all resources on process termination, not thread)
- Pointer is returned after thread join – use e.g. for returning result



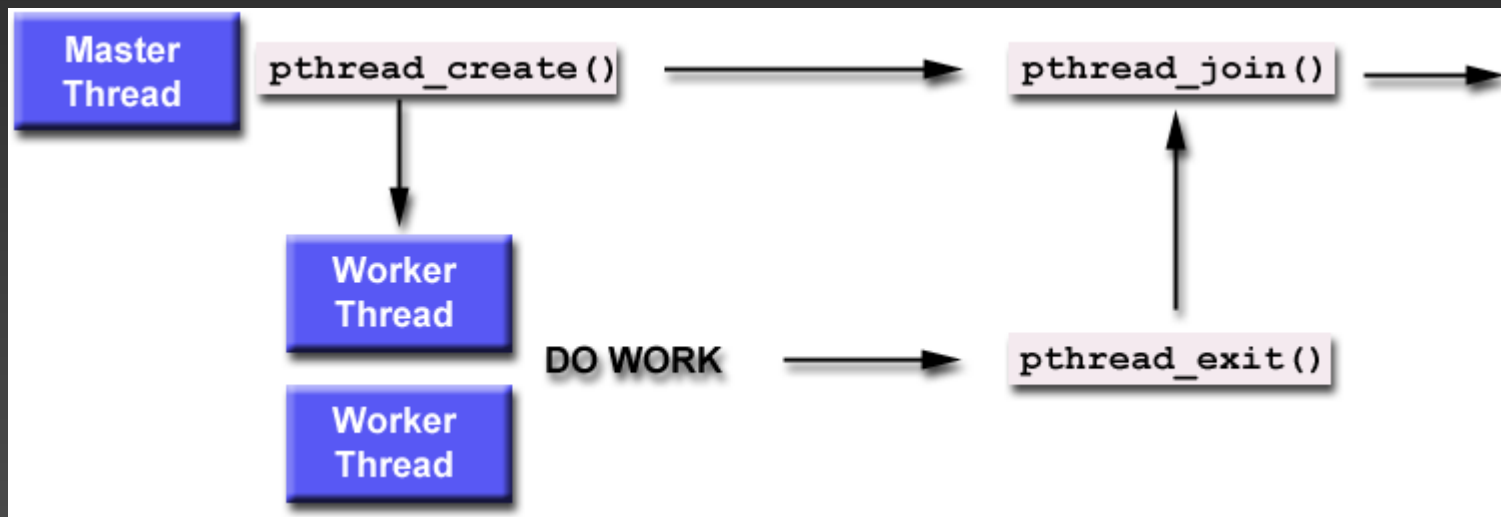
# Thread cancelling

- `int pthread_cancel (pthread_t *thread_handle)`
  - Request for `thread_handle` thread termination
  - Addressed thread may or may not terminate (just request)
  - Thread may cancel itself
  - Cancel request offers opportunity to do clean-up mem/file/etc. related to thread
  - Function exits after sending request – non-blocking
  - Return code 0 means addressed thread exists, not that it was/is/will be terminated

# Thread Join

```
int pthread_join (pthread_t thread_handle,  
void **ptr_value);
```

- blocking wait for thread thread\_handle finish
- Value ptr\_value is pointer to pointer, specified in thread\_handle in pthread\_exit()
- Necessary if we want to know exit status code



# Mutex

- Init

```
pthread_mutex_t mutex;
```

```
int pthread_mutex_init(pthread_mutex_t *mutex, const  
pthread_mutexattr_t *mutexattr);
```

- NULL attr = default

- or macro

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
```

- Usage

```
int pthread_mutex_lock(pthread_mutex_t *mutex);
```

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);
```

```
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

# Sample code

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5

void *PrintHello(void *threadid)
{ printf("%d: Hello World!\n", threadid);
  pthread_exit(NULL);
}

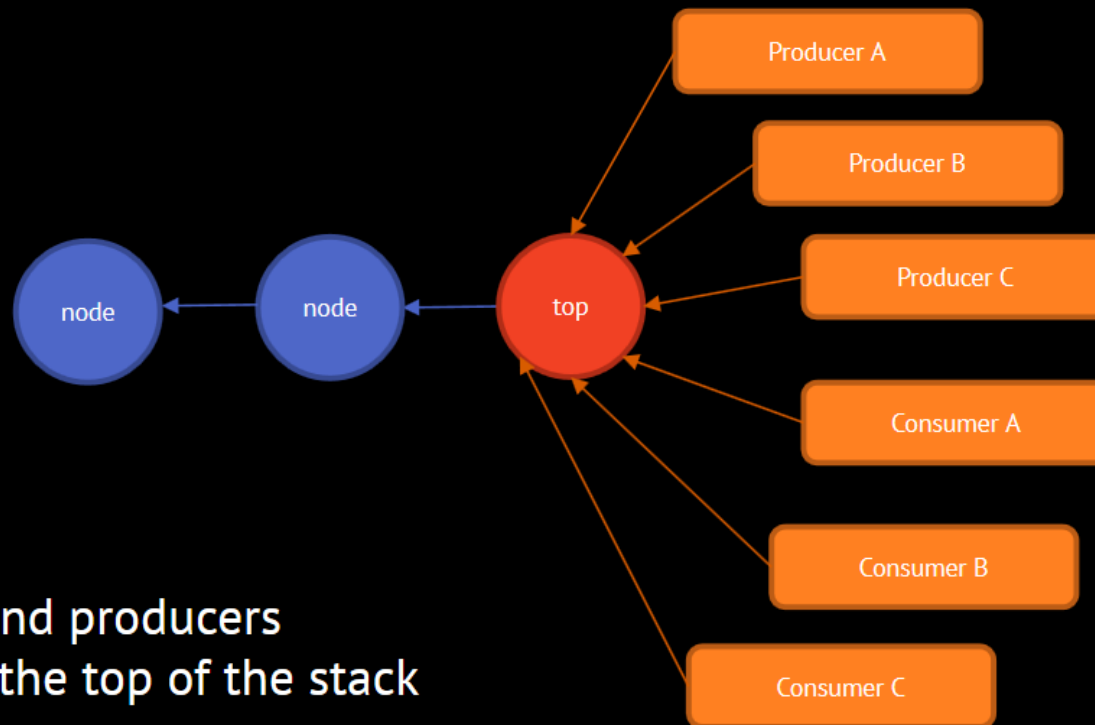
int main (int argc, char *argv[])
{ pthread_t threads[NUM_THREADS];
  for(int t=0; t<NUM_THREADS; t++)
    pthread_create(&threads[t], NULL, PrintHello, (void *)t);
  pthread_exit(NULL);
}
```

# Producer – Consumer

- Thread P produces data, C consumes
- Possible solution
  - shared data storage + mutex
- Better
  - common queue
  - counting semaphore, P increases, C decreases
  - at zero C can be put to sleep
  - problem: more C or P – removing and inserting is not atomic, PxC resource overwriting
- Best
  - conditional variables
  - any amount of C a P, single storage, no busy wait

# Data Structures


# ConcurrentStack



Consumers and producers  
compete for the top of the stack

# Data Structures


## Concurrent Queue



Consumers compete for the head.  
Producers compete for the tail.  
Both compete for the same node if the queue is empty.

# Data Structures

## ConcurrentBag





# Producer - Consumer

- Slightly better (one P, more C)
  - more queues with pointers: P and each C has its own Q, one is common
  - single mutex for pointer switching
  - if C empty, lock common, switch for its own, release mutex
    - beware, busy-wait, live-lock!
  - P after each piece of data produced: lock mutex, switch common with its own, release

# HW Vlákna v C++

- knihovna <thread>
- kolik máme v CPU hw vláken? (včetně SMT)

```
#include <iostream>
#include <thread>
#include <chrono>

int main(int argc, char** argv) {
    const unsigned int hw = std::thread::hardware_concurrency();

    // Do I care, what data type it is? NO! Auto-deduce type...
    auto hw2 = std::thread::hardware_concurrency();

    std::cout << "Counting threads...\n";           // End-Of-Line without flush (faster)
    std::this_thread::sleep_for(std::chrono::milliseconds(25));
    std::this_thread::sleep_for(std::chrono::seconds(3));
    std::cout << "Got HW threads: " << hw << std::endl; // End-Of-Line with implicit flush (safer, slower)

    return EXIT_SUCCESS;
}
```

# Fork & join

- vytvoření vlákna vs. skupiny vláken, předání parametru

```
#include <iostream>
#include <thread>

void thread_code(void) {
    std::cout << "Hello world from thread: " <<
std::this_thread::get_id() << "\n";
}

int main(int argc, char** argv) {
    std::thread my_thread(thread_code);

    my_thread.join();

    return EXIT_SUCCESS;
}
```

```
#include <iostream>
#include <vector>
#include <thread>

static const int num_threads = 10;

void thread_code(const int tid) {
    std::cout << tid << std::endl;
}

int main(int argc, char** argv) {
    std::vector<std::thread> threads;

    threads.resize(num_threads);
    for (int i = 0; i < num_threads; ++i) {
        threads[i] = std::thread(thread_code, i);
    }

    for (int i = 0; i < 10; ++i) {
        threads[i].join();
    }

    return EXIT_SUCCESS;
}
```

# Rozlišení vláken

- zjistíme hlavní vlákno, podle toho rozhodneme co dělat

```
#include <iostream>
#include <thread>

std::thread::id main_thread_id = std::this_thread::get_id();

void am_i_main(void)
{
    if (main_thread_id == std::this_thread::get_id())
        std::cout << "This is the main thread.\n";
    else
        std::cout << "This is not the main thread.\n";
}

void thread_code(void) {
    std::cout << "Hello world from thread: " << std::this_thread::get_id() << ". ";
    am_i_main();
}

int main(int argc, char** argv) {
    std::thread my_thread(thread_code);

    am_i_main();

    my_thread.join();

    return EXIT_SUCCESS;
}
```

```
This is the main thread.
Hello world from thread: 3832. This is not the main thread.
```

# Atomic

- synchronizovaná komunikace mezi vlákny
  - vhodné jen pro malé datové objemy

```
#include <iostream>
#include <vector>
#include <thread>
#include <atomic>

static const int num_threads = 10;

void thread_code(const int tid, std::atomic<int>& result) {
    std::cout << tid << std::endl;

    result += 1;
}

int main(int argc, char** argv) {
    std::vector<std::thread> threads;
    std::atomic<int> result(0);

    threads.resize(num_threads);
    for (int i = 0; i < num_threads; ++i) {
        threads[i] = std::thread(thread_code, i, std::ref(result));
    }

    for (int i = 0; i < 10; ++i) {
        threads[i].join();
    }

    std::cout << "Result: " << result << std::endl;

    return EXIT_SUCCESS;
}
```

# Mutex

- synchronizovaná komunikace mezi vlákny
  - v kritické sekci jsou možné i složitější operace

```
#include <iostream>
#include <vector>
#include <chrono>
#include <thread>
#include <mutex>

static std::mutex my_mutex;

static const int num_threads = 10;

void thread_code(const int tid, int& result) {
    std::this_thread::sleep_for(std::chrono::seconds(1));
    std::thread::id this_id = std::this_thread::get_id();

    // try to move lock BELOW printing
    my_mutex.lock();
    std::cout << "I am " << tid << " with id " << this_id << std::endl;

    result += 1;
    my_mutex.unlock();
}

int main(int argc, char** argv) {
    std::vector<std::thread> threads;
    int result = 0;

    threads.resize(num_threads);
    for (int i = 0; i < num_threads; ++i) {
        threads[i] = std::thread(thread_code, i, std::ref(result));
    }

    for (int i = 0; i < num_threads; ++i) {
        threads[i].join();
    }

    std::cout << "Result: " << result << std::endl;

    return EXIT_SUCCESS;
}
```

# • Podmíněné proměnné umožňují usínání a probouzení vláken

```
#include <iostream>,#include <string>,#include <thread>,#include  
<mutex>,#include <condition_variable>
```

```
std::mutex m;  
std::condition_variable cv;  
std::string data;  
bool ready = false;  
bool processed = false;
```

```
void worker_thread()  
{  
    // Wait until main() sends data  
    std::unique_lock<std::mutex> lk(m);  
    cv.wait(lk, [] {return ready; });  
  
    // after the wait, we own the lock.  
    std::cout << "Worker thread is processing data\n";  
    data += " after processing";  
  
    // Send data back to main()  
    processed = true;  
    std::cout << "Worker thr. signals data processing completed\n";  
  
    // Manual unlocking is done before notifying, to avoid waking up  
    // the waiting thread only to block again (see notify_one for  
    details)  
    lk.unlock();  
    cv.notify_one();  
}
```

```
int main()  
{  
    std::thread worker(worker_thread);  
  
    data = "Example data";  
    // send data to the worker thread  
    {  
        std::lock_guard<std::mutex> lk(m);  
        ready = true;  
        std::cout << "main() signals data ready for processing\n";  
    }  
    cv.notify_one();  
  
    // wait for the worker  
    {  
        std::unique_lock<std::mutex> lk(m);  
        cv.wait(lk, [] {return processed; });  
    }  
    std::cout << "Back in main(), data = " << data << '\n';  
    worker.join();  
}
```

# Úprava kontejneru na thread-safe

- přidání zámku pro vynucení exkluzivního přístupu

```
#include <deque>
#include <iostream>           // std::cout
#include <mutex>               // std::mutex, std::scoped_lock
#include <condition_variable> // std::condition_variable

template<typename T>
class synced_deque {
protected:
    std::mutex mux;
    std::deque<T> de_queue;
    std::condition_variable cv_sleep;
    std::mutex mux_sleep;

public:
    synced_deque() = default;
    synced_deque(const synced_deque<T>&) = delete;
    virtual ~synced_deque() {
        clear();
    }

    // Returns and maintains item at front of Queue
    const T& front() {
        std::scoped_lock lock(mux);
        return de_queue.front();
    }

    // Removes and returns item from front of Queue
    T pop_front() {
        std::scoped_lock lock(mux);
        auto t = std::move(de_queue.front());
        de_queue.pop_front();
        return t;
    }
}
```

```
// Adds an item to back of Queue
void push_back(const T& item) {
    std::scoped_lock lock(mux);
    de_queue.emplace_back(std::move(item));

    std::unique_lock<std::mutex> ul(mux_sleep);
    cv_sleep.notify_one();
}

// Returns true if Queue has no items
bool empty() {
    std::scoped_lock lock(mux);
    return de_queue.empty();
}

void wait() {
    while (empty()) {
        std::unique_lock<std::mutex> ul(mux_sleep);
        cv_sleep.wait(ul);
    }

    // const T& back();
    // T pop_back()
    // void push_front()
    // const T& at();
    // size_t size();
    // ...
}
```